

soberbeing78@gmail.com

+91-7021493093

Mumbai, India

[Linkedin](#)

[Github](#)

[Portfolio Site](#)

Education

B.Sc. Computer Science

2024-2027

D.G. Ruparel College, Mumbai

Current GPA: 8.00/10.00

Relevant Coursework: Data Structures, Mathematics for CS, Statistics, Python Programming, DBMS.

TECHNICAL SKILLS

Languages: Python (Advanced), SQL

Core AI & ML: TensorFlow, Keras, Scikit-learn, NumPy, Pandas, XGBoost, SHAP (Explainable AI), PyPortfolioOpt

GenAI & Agents: CrewAI, LangChain, Hugging Face (Transformers), Sentence-Transformers, Gemini APIs

Deployment & Web: FastAPI, Uvicorn, Streamlit, Git/GitHub, Git LFS, Docker, Docker-Compose

Concepts: RAG (Custom), Multi-Agent Orchestration, Vector Embeddings, Cosine Similarity, REST APIs, Prompt Engineering (CoT/ReAct), Credit Risk Modeling, RBI Compliance, Time-Series Analysis, Microservices

PROJECTS

Capital Allocation Engine | RegTech System

XGBoost, SHAP, FastAPI, Docker

- Architected a production-grade credit risk engine achieving **~88% Recall** for NPA prevention, utilizing a **parsimonious 5-feature XGBoost pipeline**.
- Deployed as a **containerized microservice (Docker)** with sub-second latency, integrating Pydantic validation to prevent drift, and leveraging **SHAP explainers** for transparent, interpretable model insights.

SubSlash AI | Autonomous Multi-Agent Marketing System

Python, CrewAI, LangChain

- Architected a 4-agent crew using **CrewAI** and **Gemini 2.5 Flash lite**, utilizing hierarchical task delegation to generate full GTM strategies.
- Implemented custom **rate-limiting logic** and **context-window management** to optimize API usage and ensure coherent long-form outputs.

DevDocs AI | RAG Q&A Bot for Developer Docs

Python, HuggingFace, RAG

- Built a RAG pipeline **from scratch without vector DBs**, implementing **Cosine Similarity search** using **NumPy** for transparent, low-latency retrieval.
- Engineered a custom orchestration script to fetch **dense embeddings** via **Hugging Face Sentence-Transformers** and inject context into the LLM prompt.

Potato Disease Classifier | End-to-End Deep Learning App

TensorFlow, CNN

- Deployed a custom **CNN architecture** (Conv2D, MaxPooling, Dropout) via a **FastAPI backend**, exposing a **REST API** endpoint for real-time inference.
- Built a prediction pipeline that handles image preprocessing/serialization and serves results to a frontend with **<100ms latency**.

CERTIFICATIONS

- Google Machine Learning Crash Course | Google (2024)